


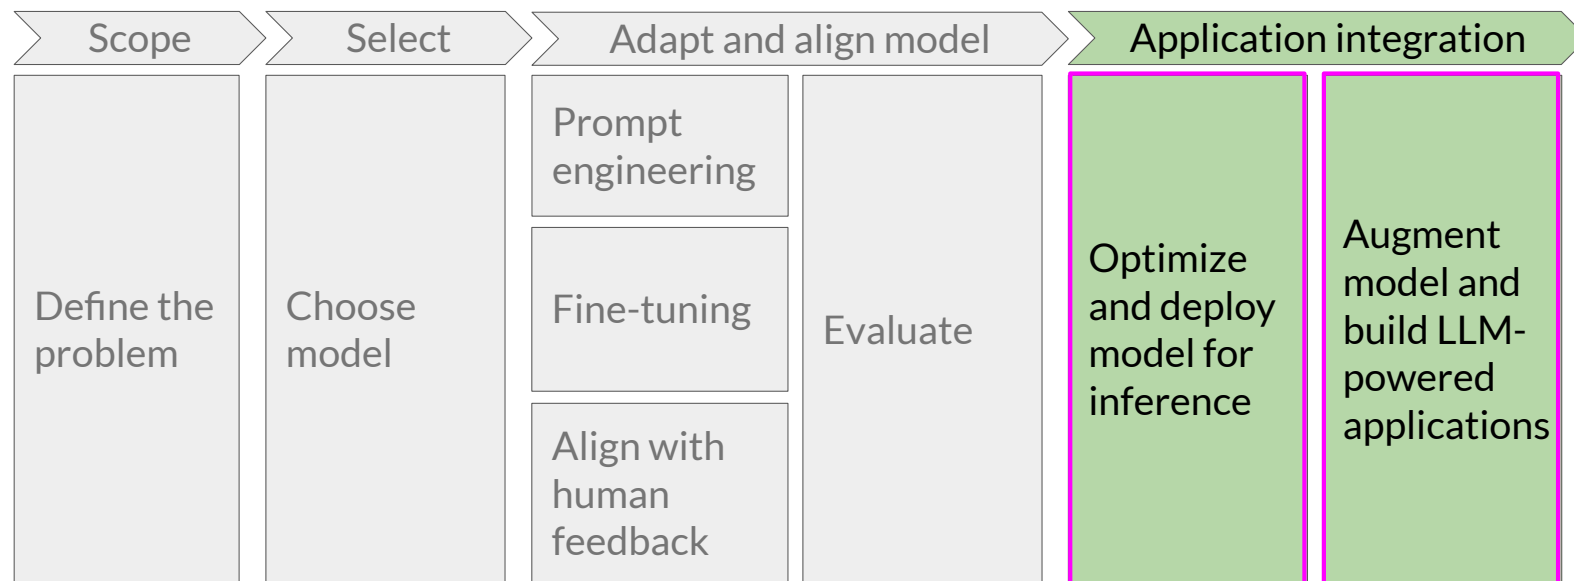
Optimize LLMs and build generative AI applications



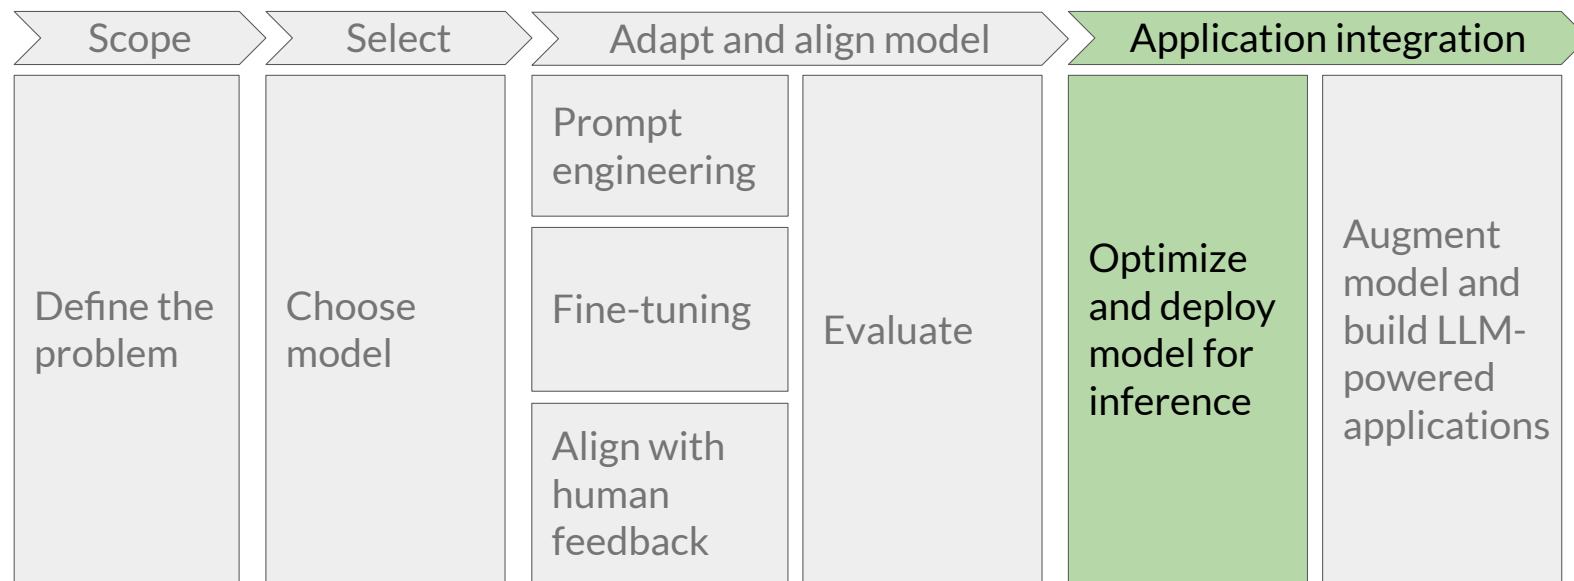
Generative AI project lifecycle



Generative AI project lifecycle




Generative AI project lifecycle




Model optimizations to improve application performance

LLM optimization techniques


Distillation



Quantization




Pruning




LLM optimization techniques


Distillation



Quantization

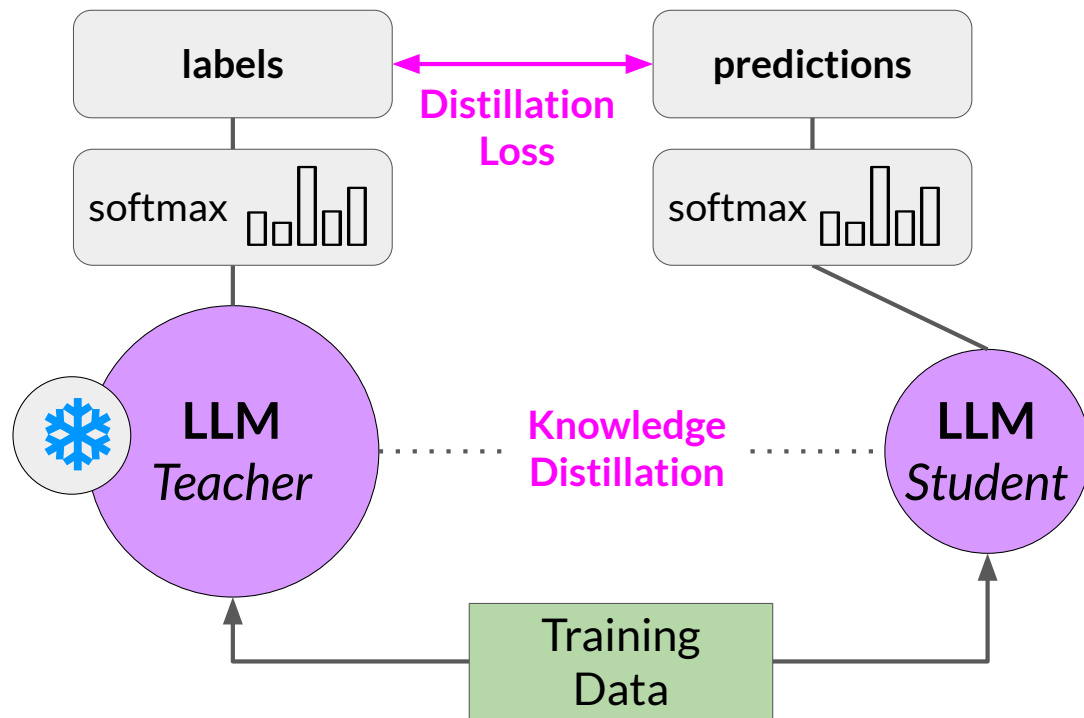


Pruning



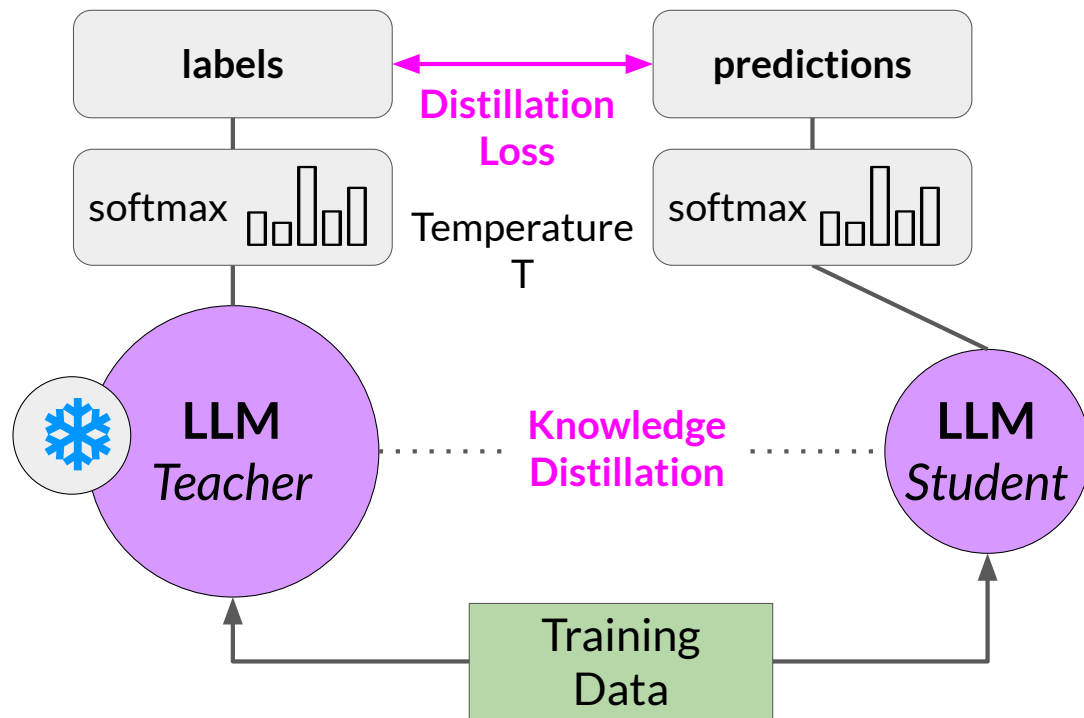
Distillation

Train a smaller student model from a larger teacher model



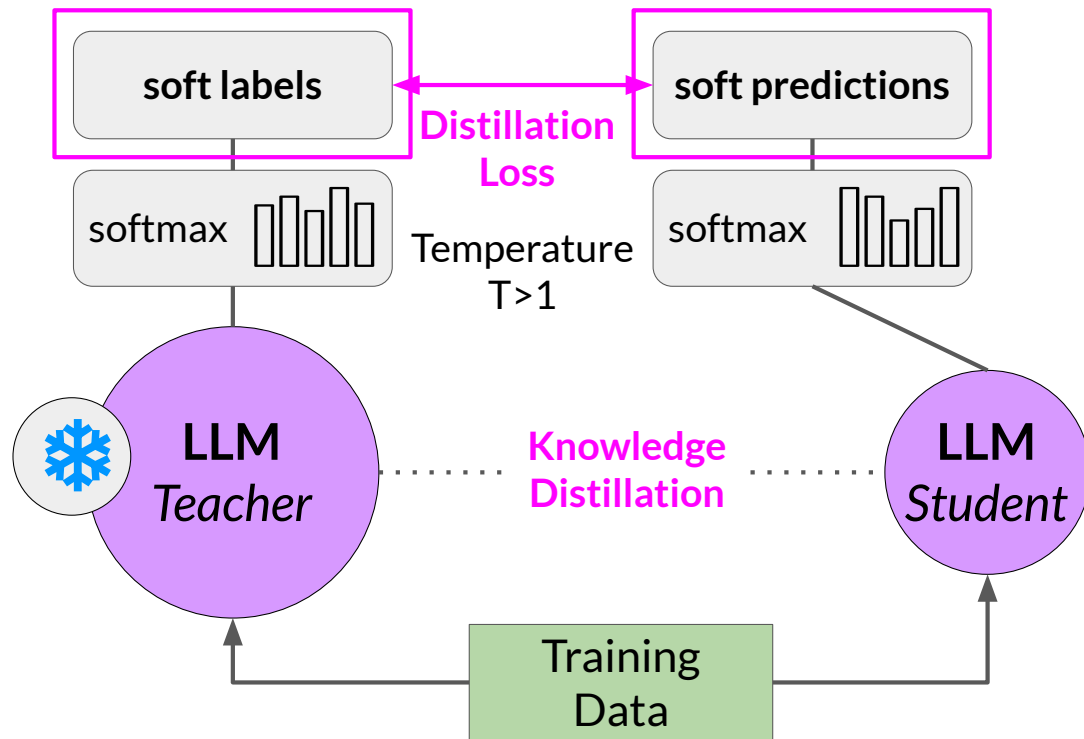
Distillation

Train a smaller student model from a larger teacher model



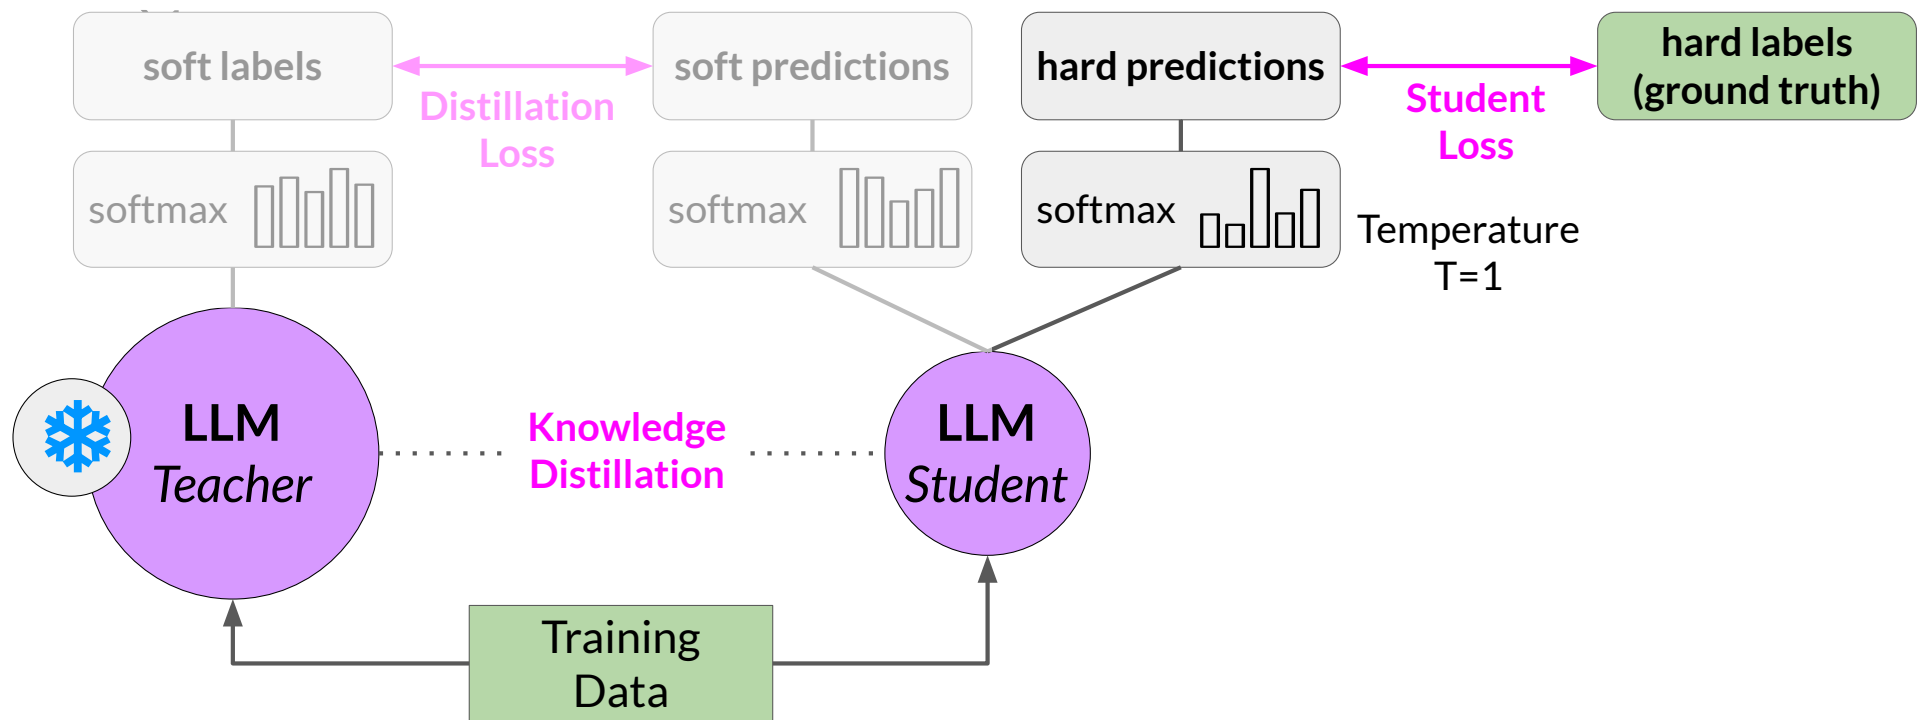
Distillation

Train a smaller student model from a larger teacher model



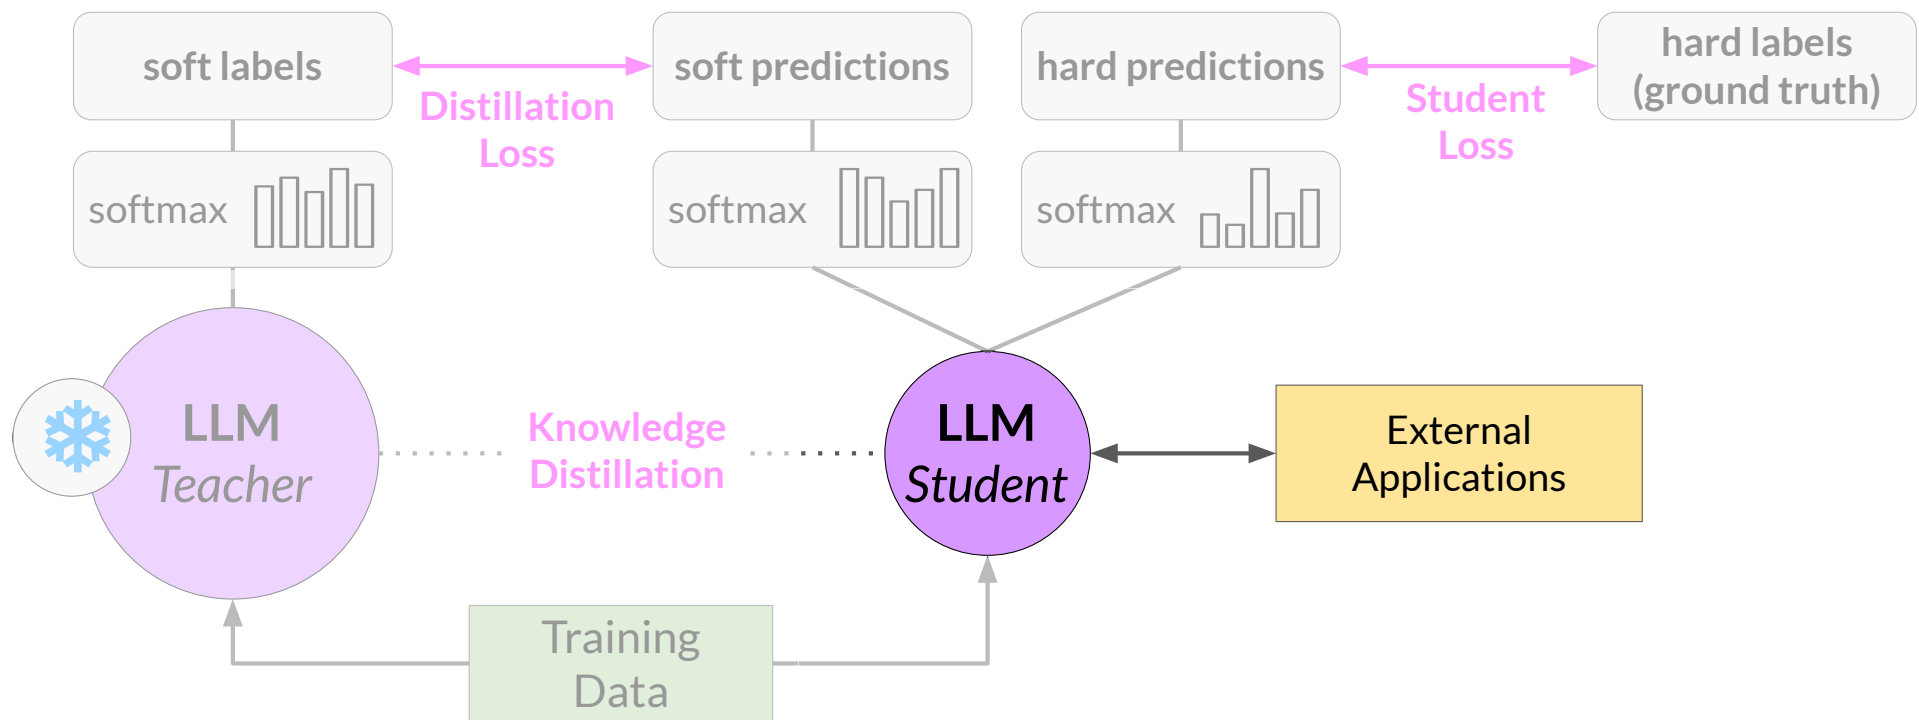
Distillation

Train a smaller student model from a larger teacher model




Distillation

Train a smaller student model from a larger teacher model




Post-Training Quantization (PTQ)

Reduce precision of model weights



Pruning

Remove model weights with values close or equal to zero



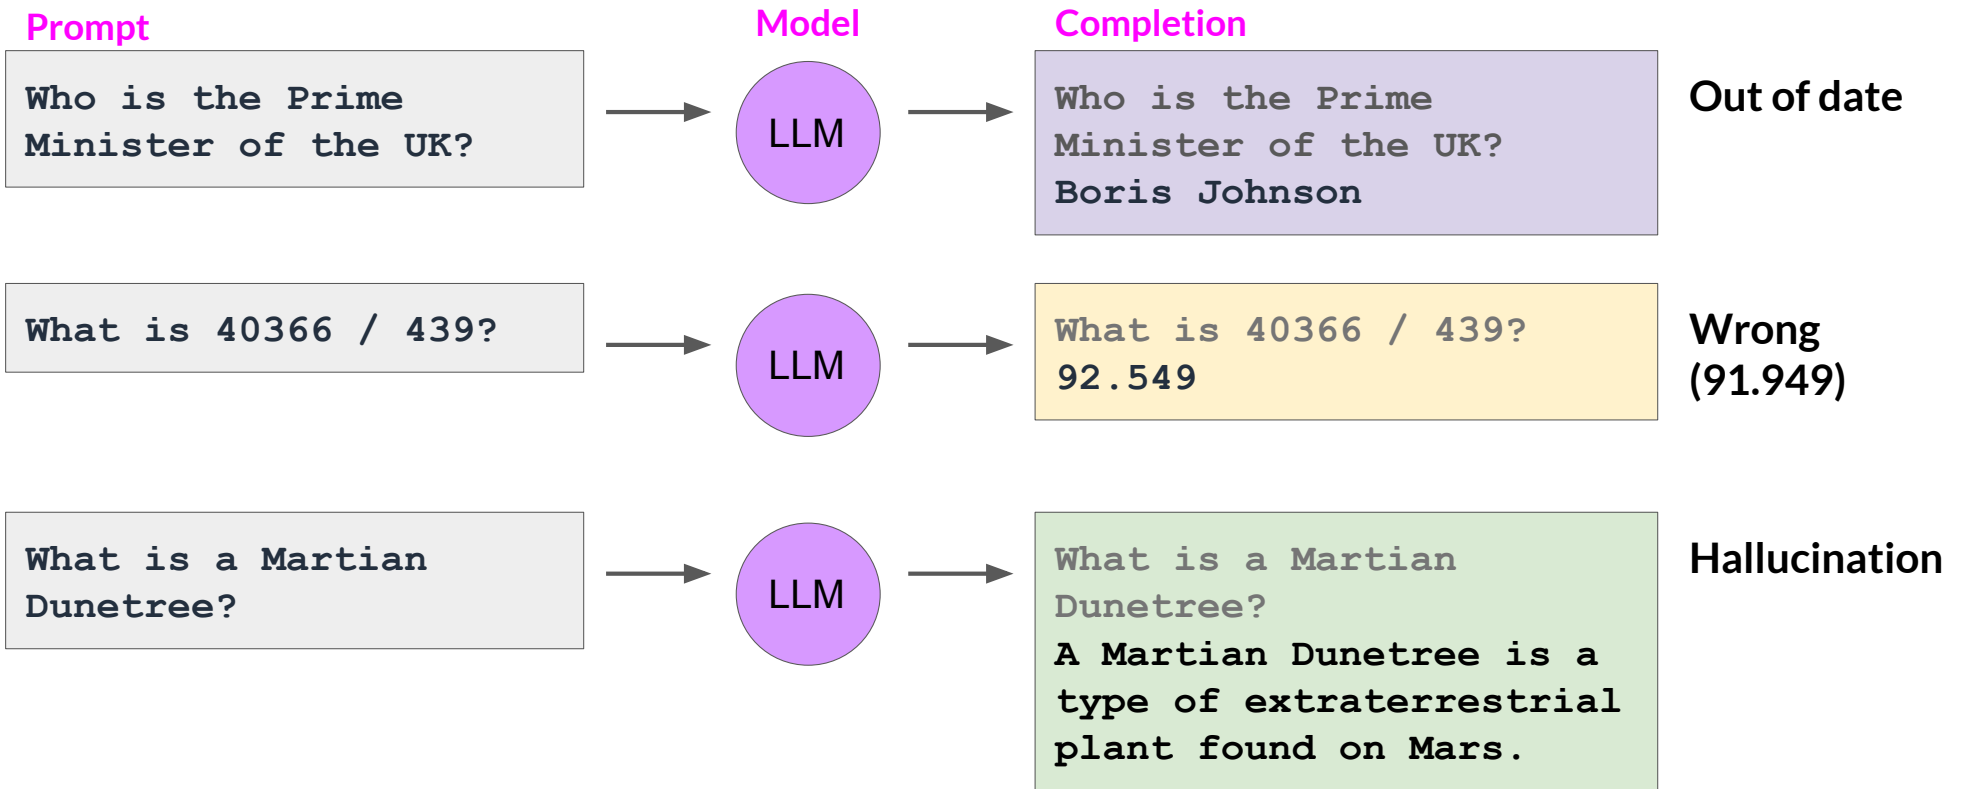
- Pruning methods
 - Full model re-training
 - PEFT/LoRA
 - Post-training
- In theory, reduces model size and improves performance
- In practice, only small % in LLMs are zero-weights

Cheat Sheet - Time and effort in the lifecycle

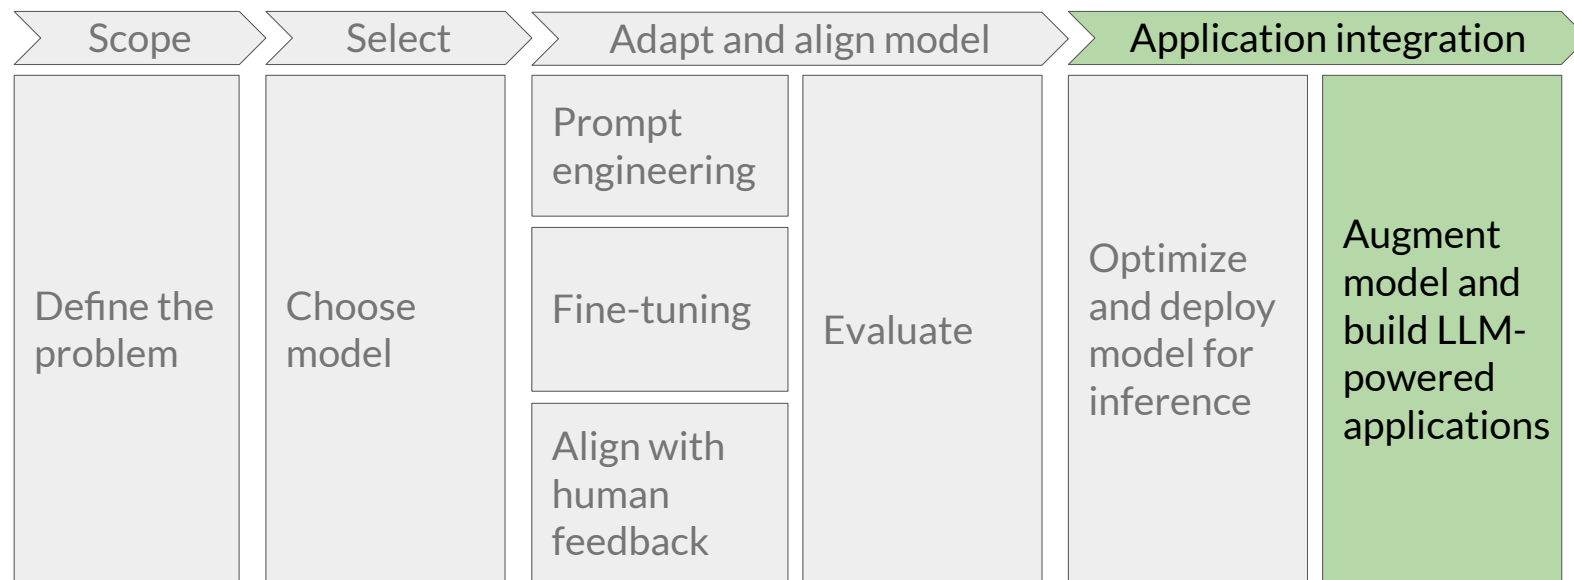
| | Pre-training | Prompt engineering | Prompt tuning and fine-tuning | Reinforcement learning/human feedback | Compression/optimization/deployment |
|-------------------|--|--|---|--|---|
| Training duration | Days to weeks to months | Not required | Minutes to hours | Minutes to hours similar to fine-tuning | Minutes to hours |
| Customization | <p>Determine model architecture, size and tokenizer.</p> <p>Choose vocabulary size and # of tokens for input/context</p> <p>Large amount of domain training data</p> | <p>No model weights</p> <p>Only prompt customization</p> | <p>Tune for specific tasks</p> <p>Add domain-specific data</p> <p>Update LLM model or adapter weights</p> | <p>Need separate reward model to align with human goals (helpful, honest, harmless)</p> <p>Update LLM model or adapter weights</p> | <p>Reduce model size through model pruning, weight quantization, distillation</p> <p>Smaller size, faster inference</p> |
| Objective | Next-token prediction | Increase task performance | Increase task performance | Increase alignment with human preferences | Increase inference performance |
| Expertise | High | Low | Medium | Medium-High | Medium |

Using the LLM in applications


Models having difficulty



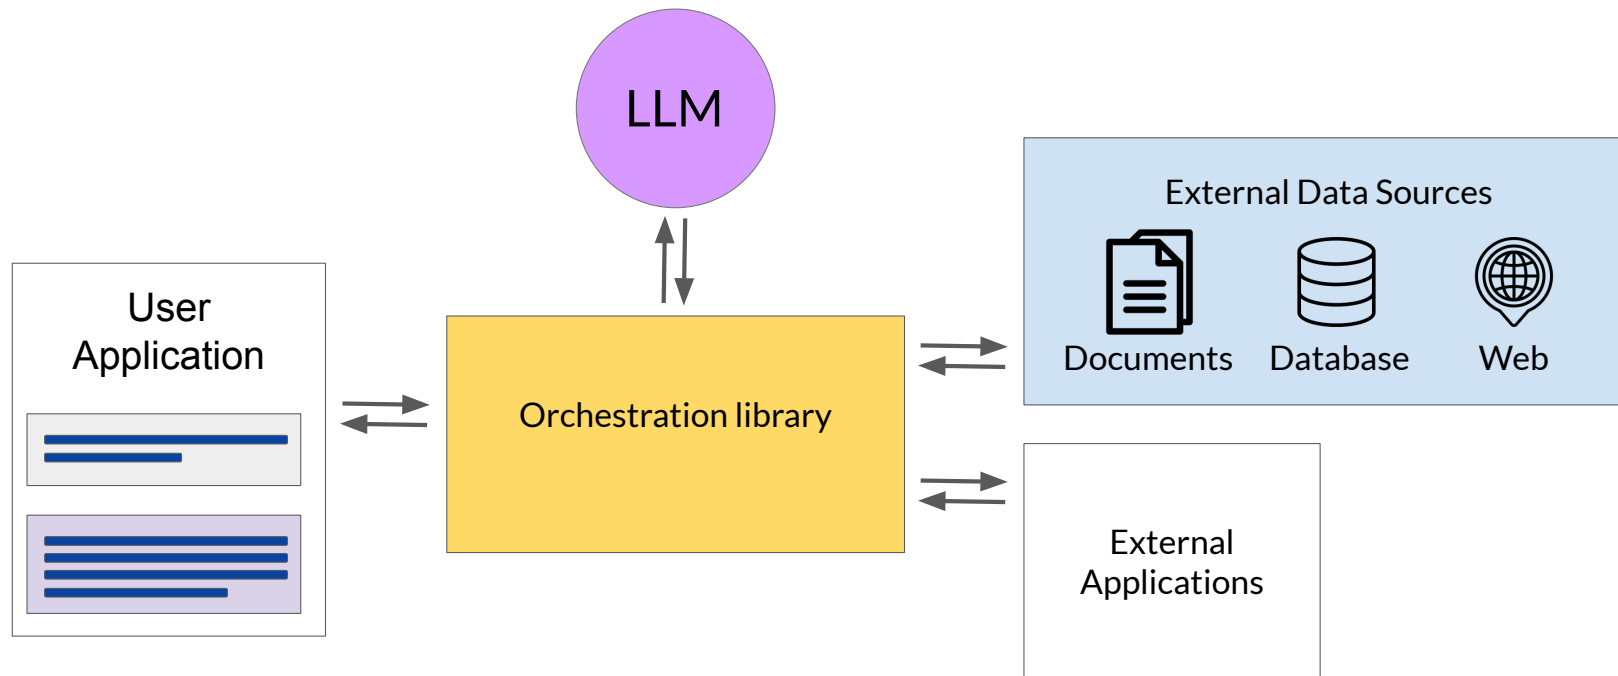
Generative AI project lifecycle



LLM-powered applications



LLM-powered applications



Retrieval augmented generation (RAG)

Knowledge cut-offs in LLMs

Prompt

Who is the
current Prime
Minister of the
United Kingdom?

Model

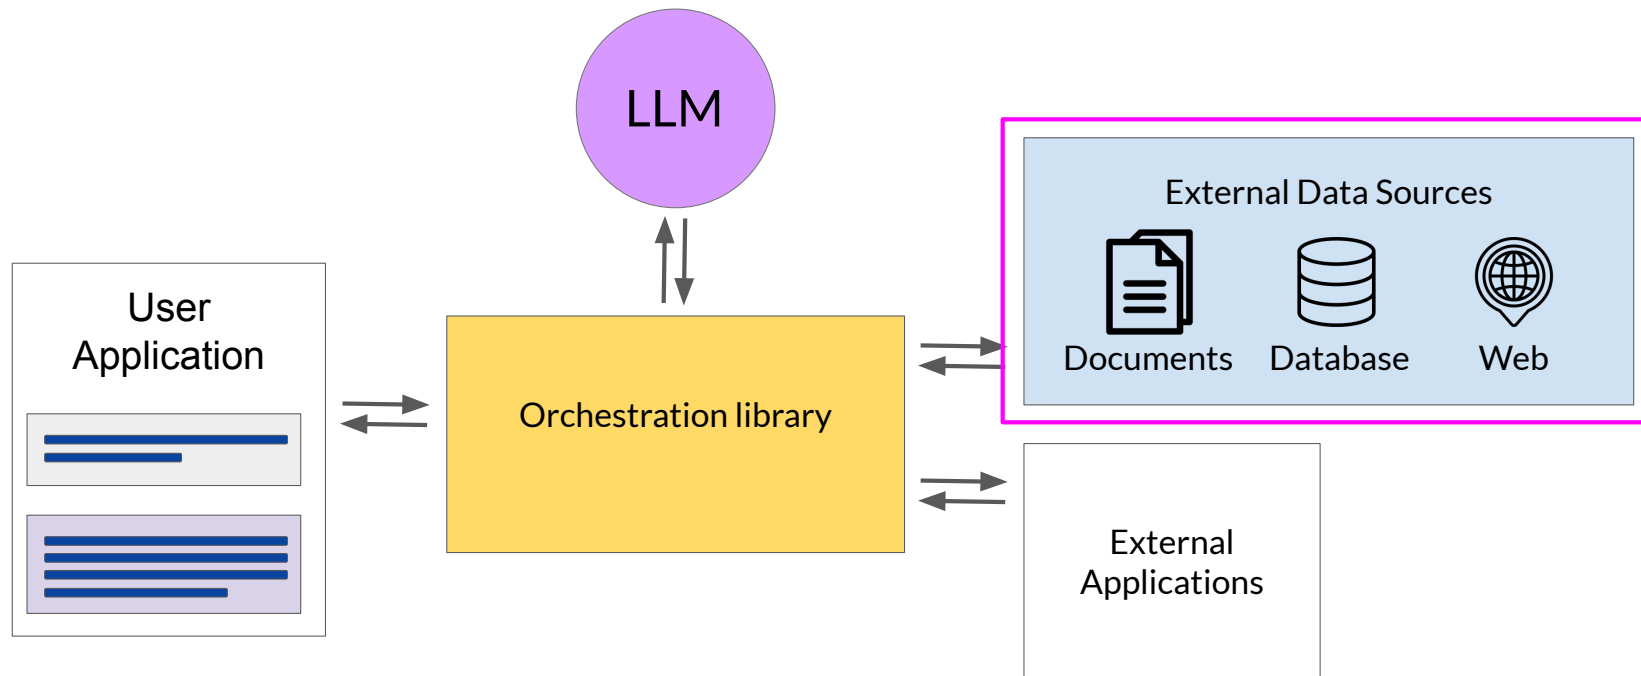
LLM

Completion

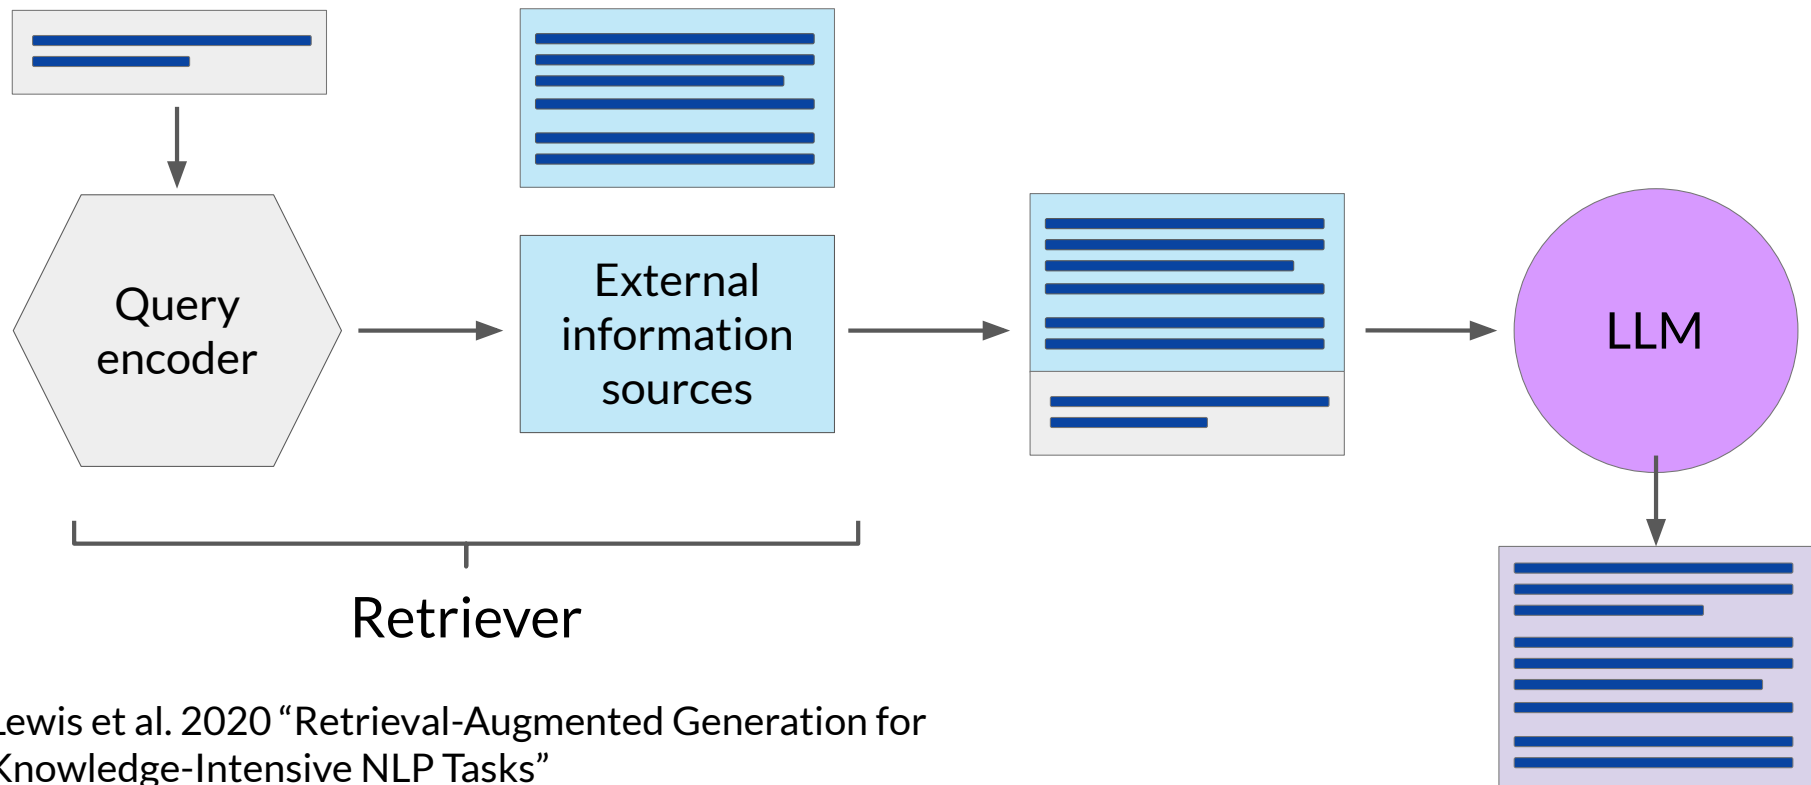
Who is the
current Prime
Minister of the
United Kingdom?

Boris Johnson

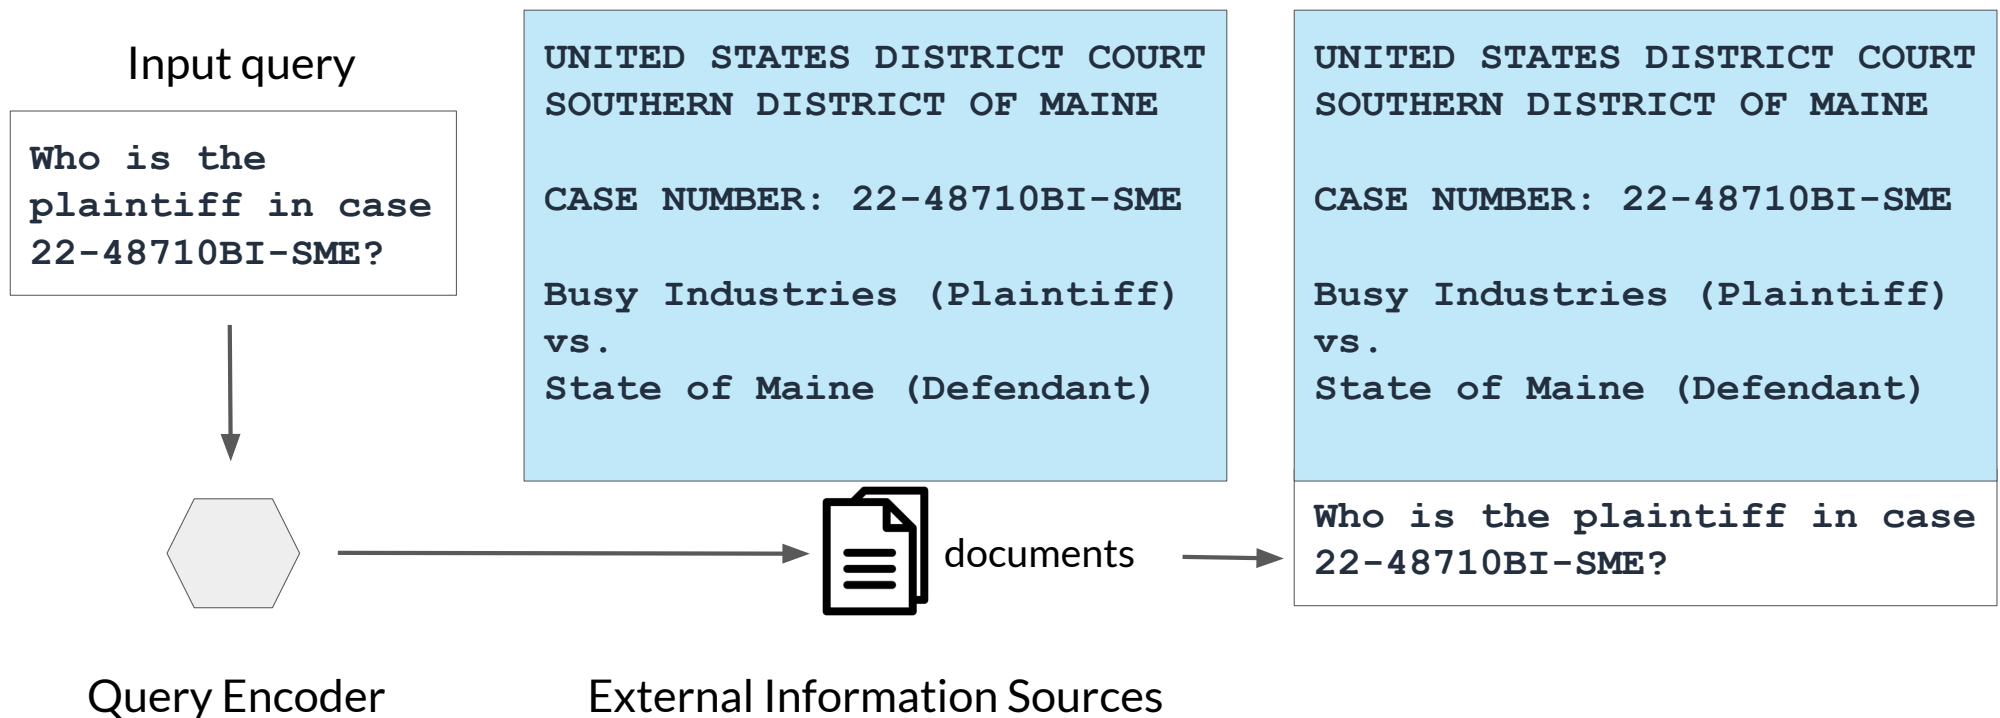
LLM-powered applications



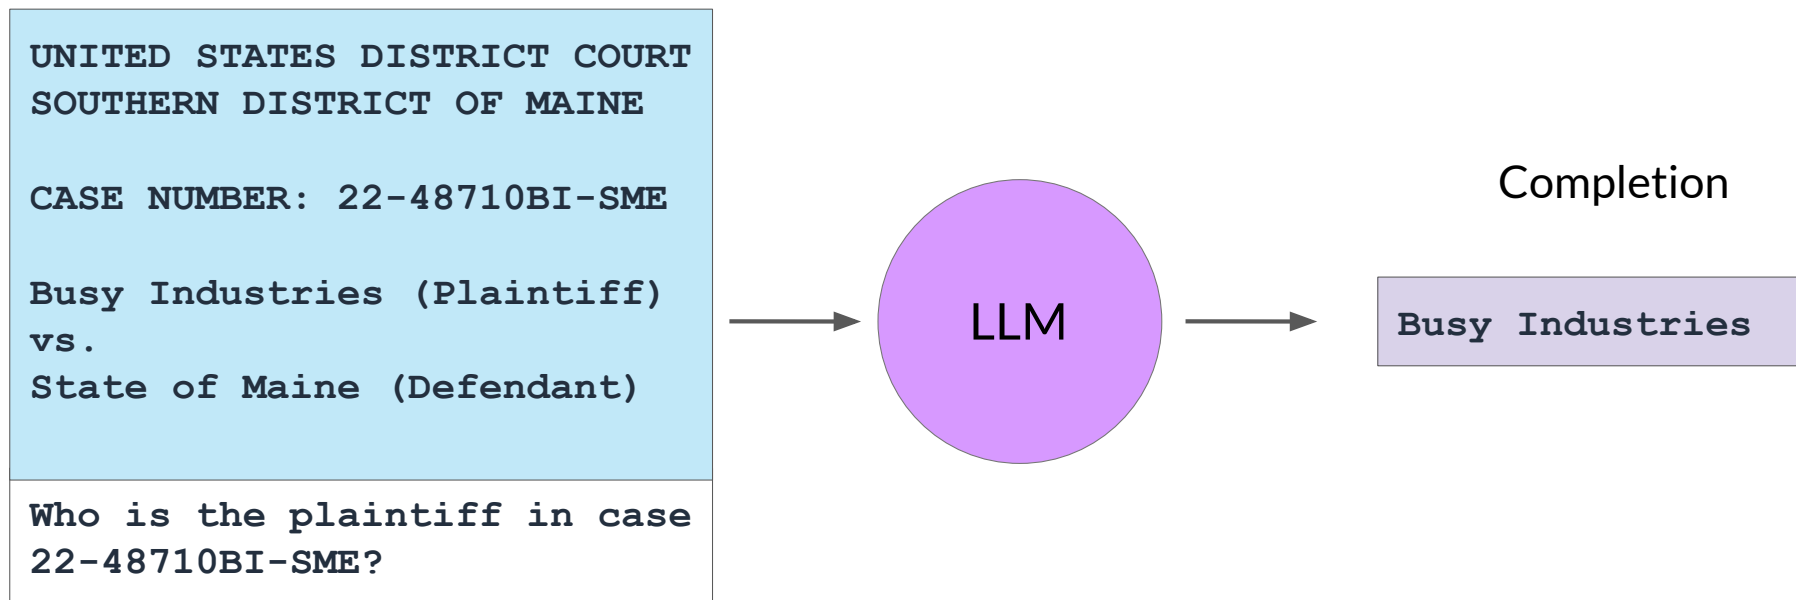
Retrieval Augmented Generation (RAG)




Example: Searching legal documents



Example: Searching legal documents



RAG integrates with many types of data sources



External Information Sources


- Documents
- Wikis
- Expert Systems
- Web pages
- Databases
- Vector Store

Data preparation for vector store for RAG

Two considerations for using external data in RAG:

1. Data must fit inside context window


Prompt context limit
few 1000 tokens



Single document too
large to fit in window



Split long sources into
short chunks




Data preparation for RAG

Two considerations for using external data in RAG:

1. Data must fit inside context window
2. Data must be in format that allows its relevance to be assessed at inference time: **Embedding vectors**

Prompt text converted to embedding vectors

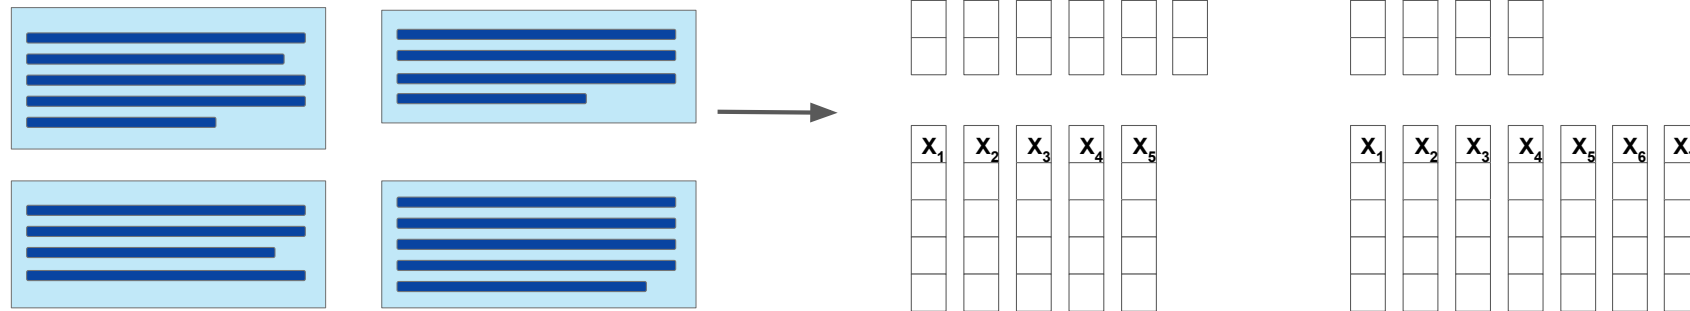


Data preparation for RAG

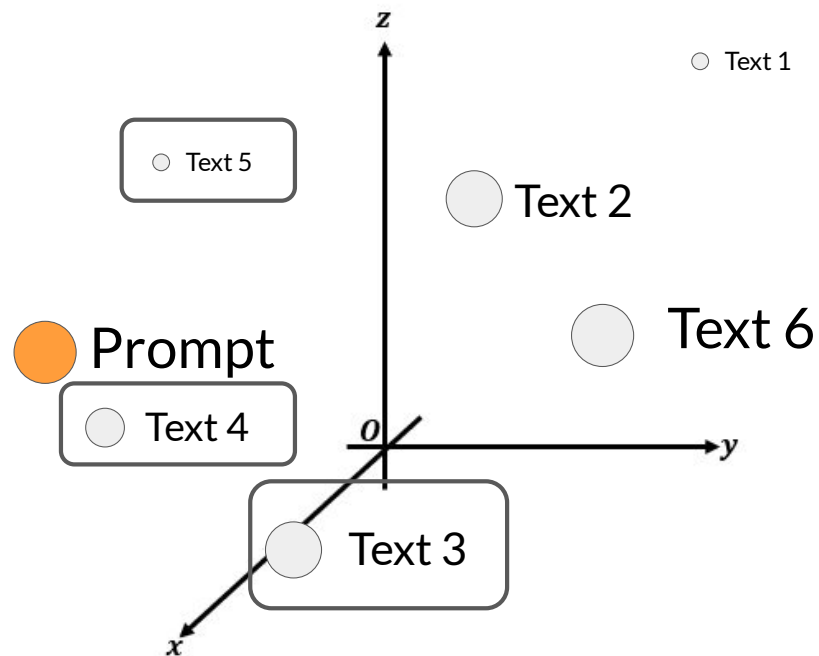
Two considerations for using external data in RAG:

1. Data must fit inside context window
2. Data must be in format that allows its relevance to be assessed at inference time: **Embedding vectors**

Process each chunk with LLM to produce embedding vectors



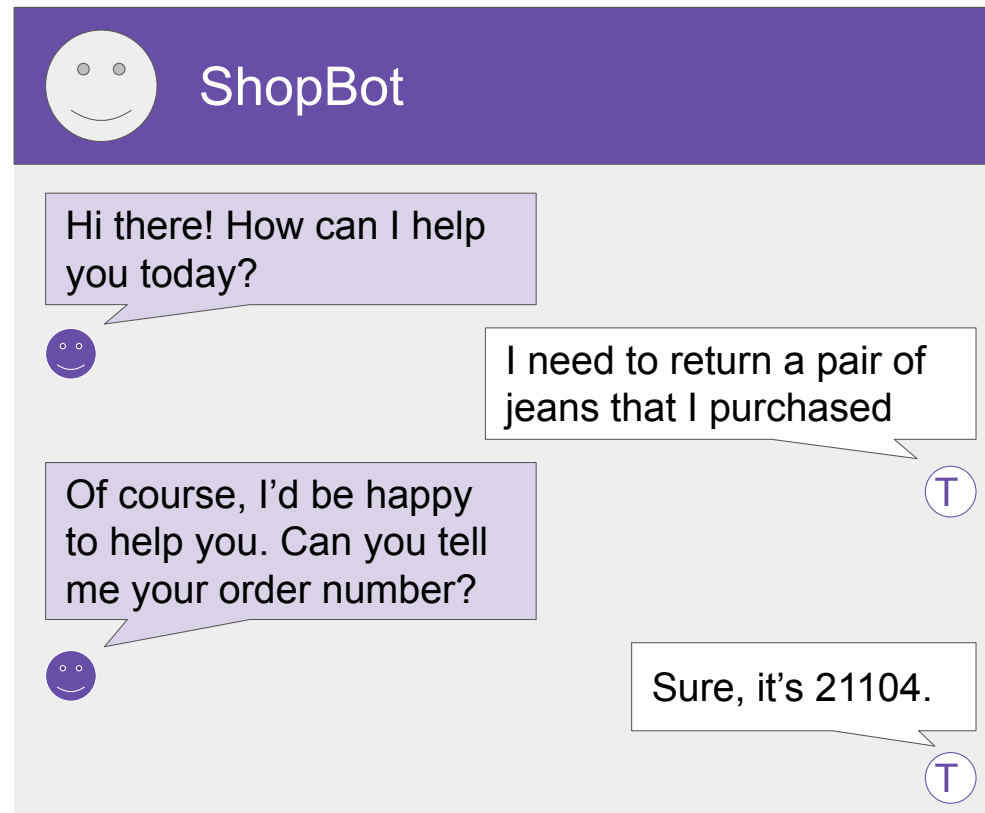
Vector database search



- Each text in vector store is identified by a key
- Enables a **citation** to be included in completion

Enabling interactions with external applications

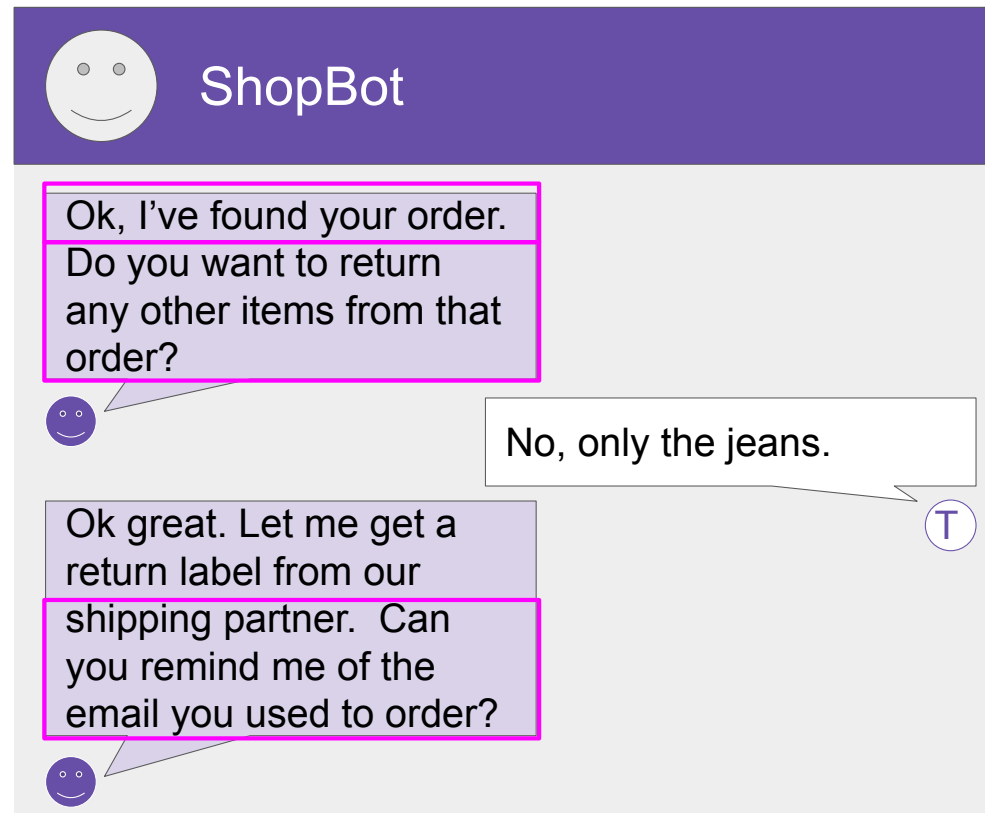
Having an LLM initiate a clothing return



Having an LLM initiate a clothing return

Lookup with RAG

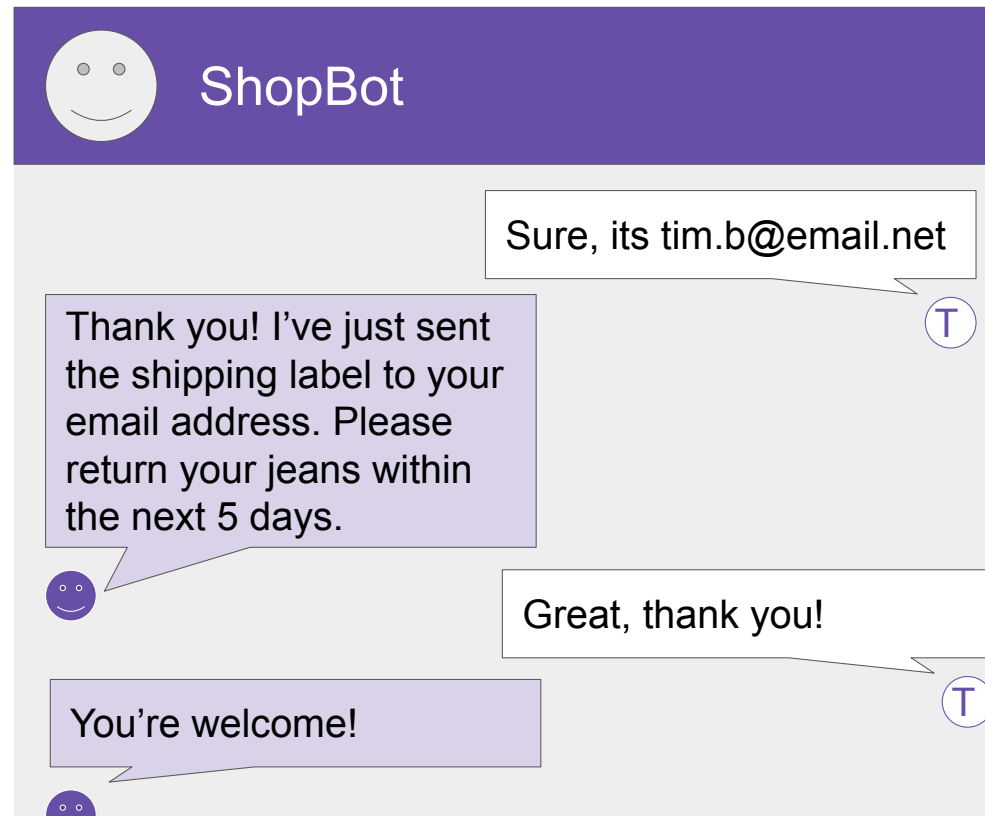
API call



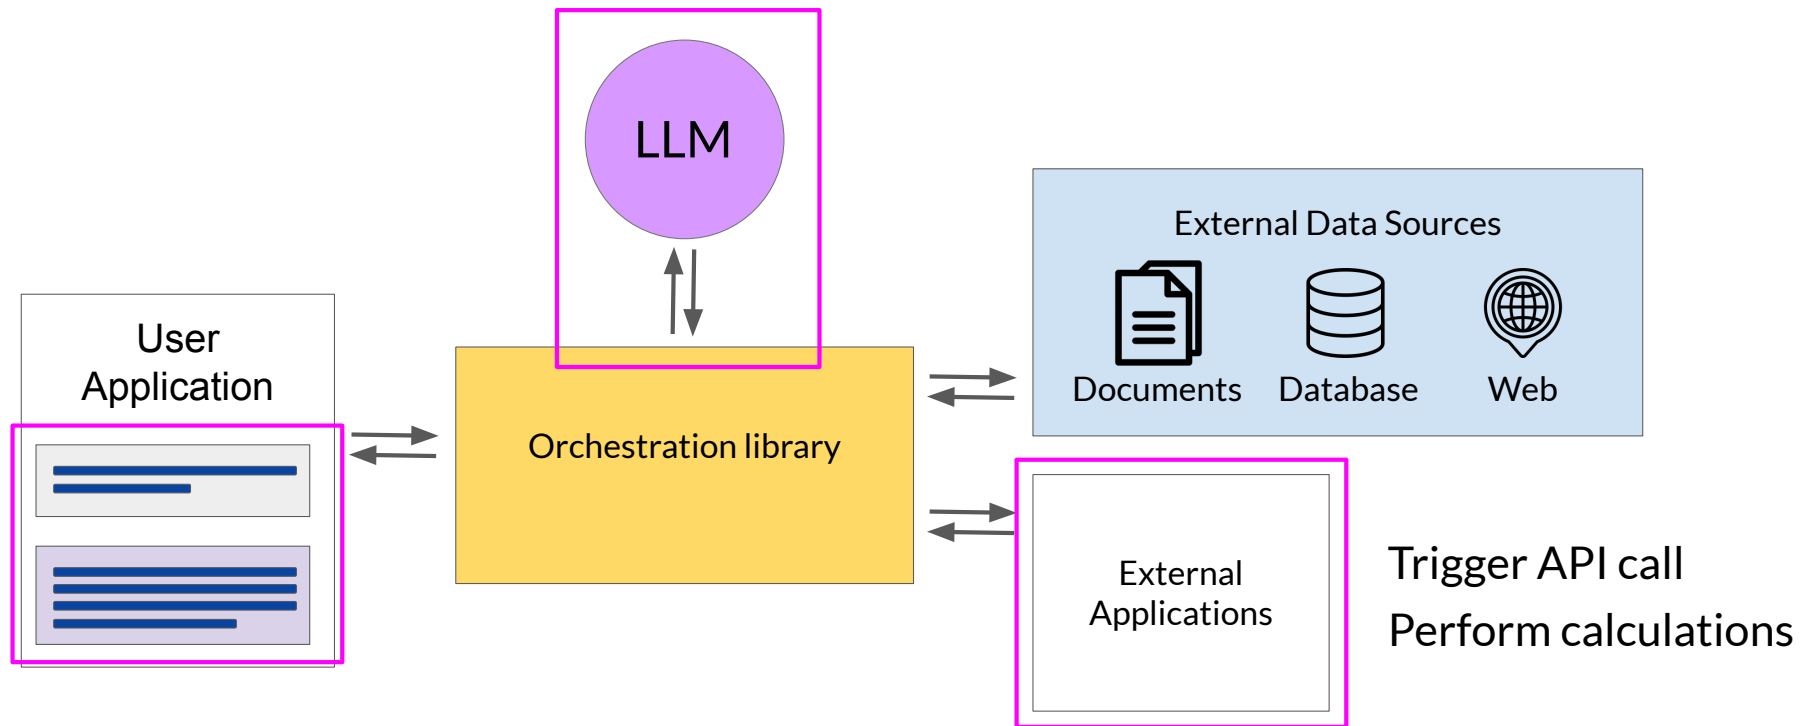
The image shows a chat interface for 'ShopBot'. At the top, there is a purple header with a smiley face icon and the text 'ShopBot'. Below this, the chat area has a light gray background. A purple message bubble from ShopBot says: 'Ok, I've found your order. Do you want to return any other items from that order?'. A white message bubble from the user responds: 'No, only the jeans.' with a blue 'T' icon. Another purple message bubble from ShopBot says: 'Ok great. Let me get a return label from our shipping partner. Can you remind me of the email you used to order?'. There are also small smiley face icons next to the user's messages.

Having an LLM initiate a clothing return

API call to the shipper



LLM-powered applications



Requirements for using LLMs to power applications

Plan actions

Steps to process return:

Step 1: Check order ID

Step 2: Request label

Step 3: Verify user email

Step 4: Email user label

Format outputs

SQL Query:

SELECT COUNT(*)

FROM orders

WHERE order_id = 21104

Validate actions

Collect required user information and make sure it is in the completion

User email:

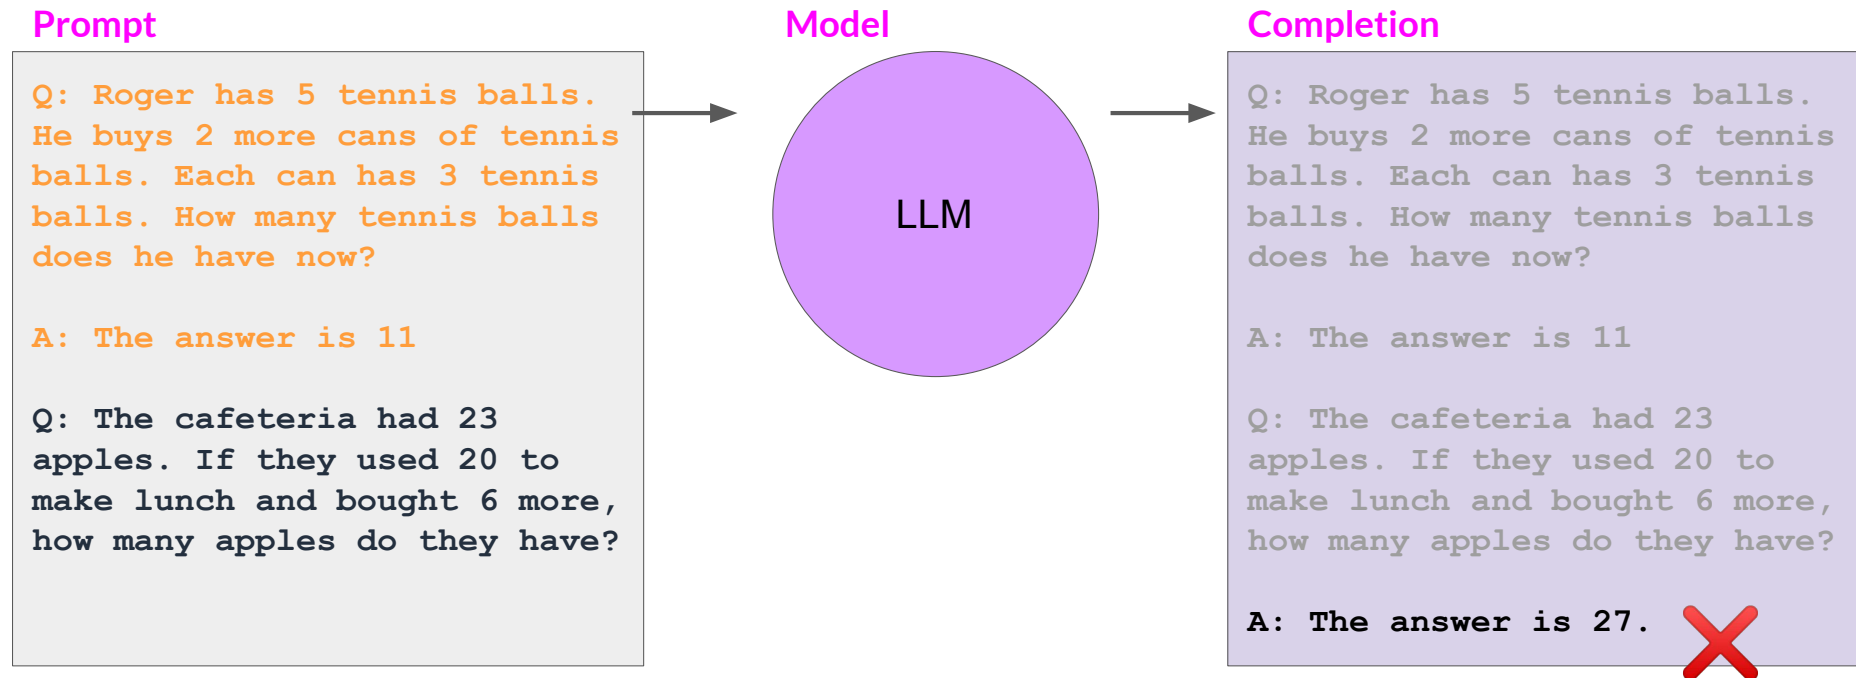
tim.b@email.net

Prompt structure is important!

```
graph TD; A[Plan actions] --> B[Format outputs]; B --> C[Validate actions]; D[Prompt structure is important!] --> A; D --> B; D --> C;
```

Helping LLMs reason and plan with Chain-of-Thought Prompting

LLMs can struggle with complex reasoning problems



Humans take a step-by-step approach to solving complex problems

Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

Start: Roger started with 5 balls.

Step 1: 2 cans of 3 tennis balls each is 6 tennis balls.

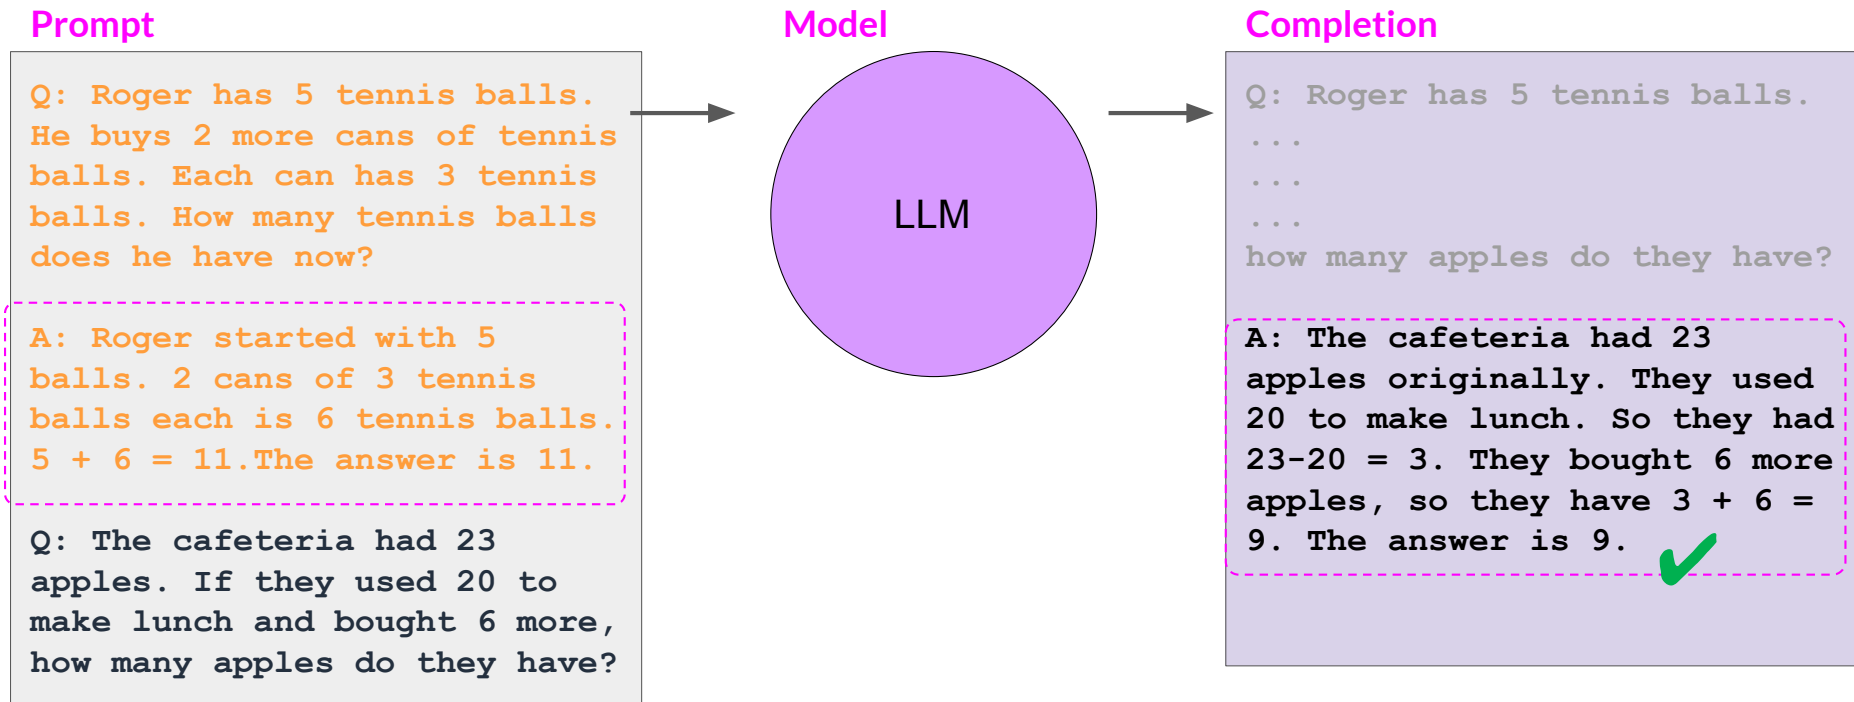
Step 2: $5 + 6 = 11$

End: The answer is 11

Reasoning steps

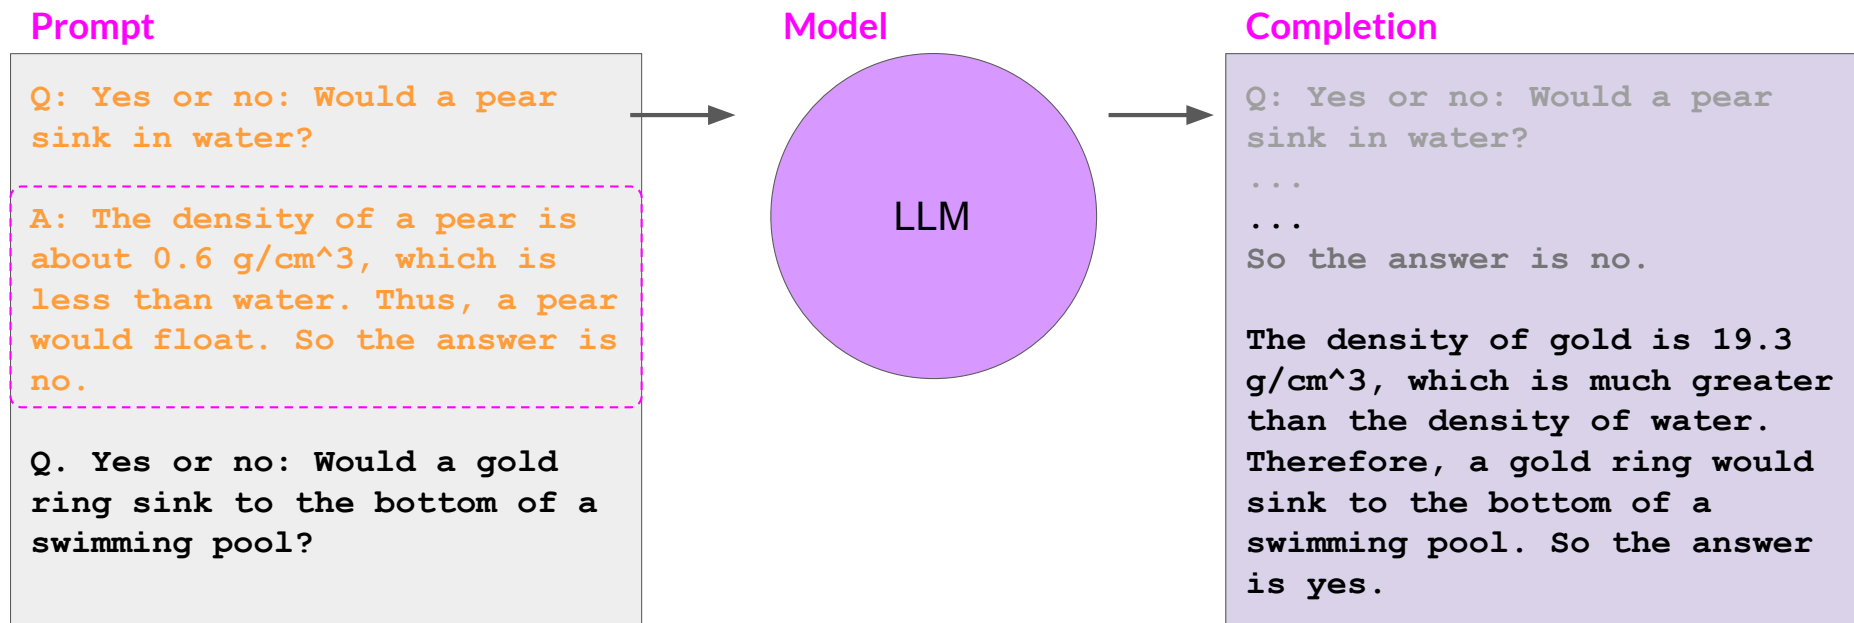
“Chain of thought”

Chain-of-Thought Prompting can help LLMs reason



Source: Wei et al. 2022, "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models"

Chain-of-Thought Prompting can help LLMs reason



Source: Wei et al. 2022, "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models"

Program-aided Language Models

LLMs can struggle with mathematics

Prompt

What is $40366 / 439$?


Model

LLM

Completion

What is $40366 / 439$?
92.549

Program-aided language (PAL) models



Chain-of-Thought (Wei et al., 2022)

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 tennis balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The bakers at the Beverly Hills Bakery baked 200 loaves of bread on Monday morning. They sold 93 loaves in the morning and 39 loaves in the afternoon. A grocery store returned 6 unsold loaves. How many loaves of bread did they have left?

Model Output

A: The bakers started with 200 loaves. They sold 93 in the morning and 39 in the afternoon. So they sold $93 + 39 = 132$ loaves. The grocery store returned 6 loaves. So they had $200 - 132 - 6 = 62$ loaves left. The answer is 62.



Program-aided Language models (this work)

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 tennis balls.

`tennis_balls = 5`

`2 cans of 3 tennis balls each is`

`bought_balls = 2 * 3`

`tennis_balls`. The answer is

`answer = tennis_balls + bought_balls`

Q: The bakers at the Beverly Hills Bakery baked 200 loaves of bread on Monday morning. They sold 93 loaves in the morning and 39 loaves in the afternoon. A grocery store returned 6 unsold loaves. How many loaves of bread did they have left?

Model Output

A: The bakers started with 200 loaves

`loaves_baked = 200`

`They sold 93 in the morning and 39 in the afternoon`

`loaves_sold_morning = 93`

`loaves_sold_afternoon = 39`

`The grocery store returned 6 loaves.`

`loaves_returned = 6`

The answer is

`answer = loaves_baked - loaves_sold_morning`
`- loaves_sold_afternoon + loaves_returned`

`>>> print(answer)`

74



Source: Gao et al. 2022, "PAL: Program-aided Language Models"

PAL example

Prompt with one-shot example

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

Answer:

```
# Roger started with 5 tennis balls  
tennis_balls = 5  
# 2 cans of tennis balls each is  
bought_balls = 2 * 3  
# tennis balls. The answer is  
answer = tennis_balls + bought_balls
```

Q: The bakers at the Beverly Hills Bakery baked 200 loaves of bread on Monday morning. They sold 93 loaves in the morning and 39 loaves in the afternoon. A grocery store returned 6 unsold loaves. How many loaves did they have left?

PAL example

Prompt with one-shot example

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

Answer:

Roger started with 5 tennis balls

tennis_balls = 5

2 cans of tennis balls each is

bought_balls = 2 * 3

tennis balls. The answer is

answer = tennis_balls + bought_balls

Q: The bakers at the Beverly Hills Bakery baked 200 loaves of bread on Monday morning. They sold 93 loaves in the morning and 39 loaves in the afternoon. A grocery store returned 6 unsold loaves. How many loaves did they have left?

Completion, CoT reasoning (blue), and PAL execution (pink)

Answer:

The bakers started with 200 loaves

loaves_baked = 200

They sold 93 in the morning and 39 in the afternoon

loaves_sold_morning = 93

loaves_sold_afternoon = 39

The grocery store returned 6 loaves.

loaves_returned = 6

The answer is

answer = loaves_baked
- loaves_sold_morning
- loaves_sold_afternoon
+ loaves_returned

Program-aided language (PAL) models

